

Solution ① → [220621] Assignment - 5
[Manish Kumar]

Logistic Regression vs Classic Perceptron ?

- So generally logistic regression preferable to use a logistic regression classifier rather than a classic perceptron because -
 - logistic regression outputs probabilities (values between 0 and 1), making it suitable for probabilistic interpretation and gradient based optimization.
 - The classic perceptron uses hard threshold (step function), which is not differentiable and does not provide probability estimates.
 - Logistic regression uses a smooth, differentiable sigmoid activation, which enables effective training using gradient descent.

Tweaking perceptron to match logistic Regression

- Replace the perceptron's step activation function with sigmoid (logistic) activation function
- Train the model using a loss function such as cross-entropy, instead of the perceptron's error-counting loss,

Solution - ② → Output Layer Design

- Spam vs. Ham.
 - + neurons: $\rightarrow 2$
 - + Activation \rightarrow sigmoid
- MNIST (10 digits)
 - + neurons $\rightarrow 10$
 - + Activation \rightarrow softmax.

Solution ③ → : Given

① Input layer: 10 neurons (passes through)
 Hidden layer: 50 neurons (ReLU)

Output layer: 3 neurons (ReLU)

ReLU activation function in hidden layer

④ Shape of input matrix \rightarrow (batch size, 10)

⑤ Shape of weight matrix $\rightarrow (w_b) = (10, 50)$
 Shape of bias vector (b_b) = 50

⑥ Shape of output layer weight matrix (w_o)

⑦ Output shape = (batch size, 3) $= (50, 3)$

⑧ Equation \Rightarrow
$$y = \text{ReLU}(x \cdot w_b + b_b) \cdot w_o + b_o$$

Solution ④ →

$$\therefore E(w) = - \sum_{n=1}^N \{ t_n \ln y_n + (1-t_n) \ln(1-y_n) \}$$

a_m = logistic sigmoid activation

$$y_n = \sigma(a_n)$$

$$\left\{ \begin{array}{l} \sigma(a_n) = 1 \\ 1+e^{-a_n} \end{array} \right.$$

$$y'_n = \sigma(a_n)(1-\sigma(a_n))$$

$$E = - \sum_{n=1}^N \{ t_n \ln \sigma(a_n) + (1-t_n) \ln(1-\sigma(a_n)) \}$$

$$E' = - \sum_{n=1}^N \left(t_n \frac{y'_n}{\sigma(a_n)} - (1-t_n) \cdot \frac{y'_n}{1-\sigma(a_n)} \right)$$

$$= - \sum_{n=1}^N \left(t_n (1-\sigma(a_n)) - (1-t_n) \sigma(a_n) \right)$$

$$= - \sum_{n=1}^N (t_n - t_n \sigma(a_n) - \sigma(a_n) + t_n \sigma(a_n))$$

$$= - \sum_{n=1}^N (t_n - \sigma(a_n))$$

$$\boxed{\frac{\partial E}{\partial a_K} = y_K - t_K}$$

Solution 5 -

2	5	-3	0
0	6	0	-4
-1	-3	0	2
5	0	0	3

A

-2	0
4	6

$$\times = ?$$

Let

$$A \times B = C = \begin{array}{|c|c|c|} \hline a_{11} & a_{12} & a_{13} \\ \hline a_{21} & a_{22} & a_{23} \\ \hline a_{31} & a_{32} & a_{33} \\ \hline \end{array}$$

$$a_{11} = -2 \times 2 + 5 \times 0 + 0 \times 4 + 6 \times 6 = 32$$

similarly

$$a_{12} = 14, \quad a_{13} = -18$$

$$a_{21} = -22, \quad a_{22} = -24, \quad a_{23} = 12$$

$$a_{31} = 22; \quad a_{32} = 6, \quad a_{33} = 18$$

∴

$$C =$$

32	14	-18
-22	-24	12
22	6	18

Solution ⑥ →

1D Strided convolutional layer with input having four units with activations (x_1, x_2, x_3, x_4) which is padded with zero to given

$(0, x_1, x_2, x_3, x_4, 0)$. β ~~row~~ filter with parameters (w_1, w_2, w_3) , stride = 2

1D activation vector

$$\text{Output} = \left(\begin{array}{l} w_2 x_1 + w_3 x_1 \\ w_1 x_2 + w_2 x_3 + w_3 x_4 \end{array} \right)$$

$$= A \times \begin{pmatrix} 0 \\ x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix}$$

$$A = \begin{bmatrix} 0 & w_2 & w_3 & 0 & 0 \\ 0 & 0 & w_1 & w_2 & w_3 \end{bmatrix}$$

upsampling convolutional

(z_1, z_2) — input layer activation

Output stride = 2

Output size = 6 = (input size - 1) \times stride +

$$= (2-1) \times 2 + 3 = 5 \quad \text{filter size}$$

\therefore zero padding gives us 6 dimensions
treating zero bias. Input elmts (z_1, z_2, z_3)

z_1 gets spread out using the filter as position

$$0 \cdot 1 \cdot \phi(z) : w_1 z_1, w_2 z_1, -w_3 z_1$$

z_2 gets placed 2 steps (because of stride = 2)
at position $2, 3, 4$

$$w_1 z_2, w_2 z_2, w_3 z_2$$

If overlapping happens, values are added

6th dimension

$$\text{output vector } y = \begin{bmatrix} w_1 z_1 \\ w_2 z_1 \\ w_3 z_1 + w_1 z_2 \\ w_2 z_2 \\ w_3 z_2 + w_1 z_3 \\ w_2 z_3 \end{bmatrix}$$

matrix form using transpose A^T

$$A = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & w_1 & w_2 & 0 & 0 \\ 0 & 0 & w_1 & w_2 & w_3 \\ 0 & 0 & 0 & 0 & w_1 & w_2 \end{bmatrix}$$

$$A^T = \begin{bmatrix} 0 & 0 & 0 \\ w_1 & 0 & 0 \\ w_2 & w_1 & 0 \\ 0 & w_2 & 0 \\ 0 & w_3 & w_1 \\ 0 & 0 & w_2 \end{bmatrix} \Rightarrow y = A^T \begin{bmatrix} z_1 \\ z_2 \end{bmatrix}$$

Solution 7)

$$\text{Input} = \begin{bmatrix} 1 & 1 & 2 & 1 \\ 2 & 1 & 2 & 0 \\ 3 & 2 & 4 & 1 \end{bmatrix}$$

$$\text{Encoder: } w = \begin{bmatrix} 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ -1 & 0 & 1 & 0 \end{bmatrix}, b = \begin{bmatrix} 0 \\ 0 \\ -1 \end{bmatrix}$$

Encoder output:

$$\text{ReLU}(w \cdot \text{Input} + b)$$

$$w \cdot \text{Input} + b = \begin{bmatrix} 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \\ -1 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} 1 & 1 & 2 & 1 \\ 2 & 1 & 2 & 0 \\ 3 & 2 & 4 & 1 \\ 1 & 0 & 1 & 2 \end{bmatrix}$$

$$+ \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ -1 & -1 & -1 & -1 \end{bmatrix}$$

$$= \text{ReLU} \left(\begin{bmatrix} 2 & 2 & 4 & 2 \\ 5 & 3 & 6 & 1 \\ 1 & 0 & 1 & -1 \end{bmatrix} \right) = \begin{bmatrix} 2 & 2 & 4 & 2 \\ 5 & 3 & 6 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

battle vector

$$= \text{ReLU} \left(\begin{bmatrix} 1 & 0 & 1 \\ -1 & 0 & 0 \end{bmatrix} \cdot \begin{bmatrix} 2 & 2 & 4 & 2 \\ 5 & 3 & 6 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix} + b \right) = \begin{bmatrix} 2 & 2 & 4 & 2 \\ 5 & 3 & 6 & 1 \\ 1 & 0 & 1 & 0 \end{bmatrix}$$

zero matrix given

$$= \text{ReLU} \left(\begin{bmatrix} 3 & 2 & 5 & 2 \\ 3 & 1 & 2 & 1 \end{bmatrix} \right) = \begin{bmatrix} 3 & 2 & 5 & 2 \\ 3 & 1 & 2 & 0 \end{bmatrix}$$

Decoded output:

$$= \text{ReLU} \left(\begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} 3 & 2 & 5 & 2 \\ 3 & 1 & 2 & 0 \end{bmatrix} \right)$$

calculated bottleneck

$$+ \begin{pmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{pmatrix})$$

bias

$$= \begin{bmatrix} 3 & 2 & 5 & 2 \\ 4 & 2 & 3 & 1 \\ 0 & 1 & 3 & 2 \end{bmatrix} \quad (\text{decoded output})$$

Reconstructed output $\Rightarrow \text{ReLU} \left(\begin{bmatrix} 1 & 0 & -1 \\ 1 & -1 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 1 \end{bmatrix} \begin{bmatrix} 3 & 2 & 5 & 2 \\ 4 & 2 & 3 & 1 \\ 0 & 1 & 3 & 2 \end{bmatrix} \right)$

$$= \begin{bmatrix} 3 & 1 & 2 & 0 \\ 0 & 0 & 2 & 1 \\ 1 & 2 & 4 & 3 \\ 1 & 0 & 3 & 0 \end{bmatrix} + \begin{pmatrix} 0 & 0 & 0 & 0 & -1 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ -3 & -3 & -3 & -3 & 1 \end{pmatrix}$$

$$\text{Reconstruction loss} = \frac{1}{m} \sum_{p=1}^m \sum_{j=1}^n (x_{pj} - \hat{x}_{pj})^2$$

$$= \frac{1}{16} (y + 0 + 0 + 1 + y + 1 + 0 + 1 + y + 0 + \\ f + 1 + 0 + 1 + 1 + 1)$$

②

$$= 1.375$$

MSE loss gradient

$$\frac{\partial L}{\partial y} = \frac{2}{n} (y_{\text{target}} - y_{\text{true}})$$

$$= \sum_p \sum_j y_{pj} - \left(\sum_i y_{ij} \right)_{\text{true}}$$

$$= \frac{2}{16} (23 - 25) = \frac{1}{8} \times 2 = \underline{\underline{0.25}}$$

Programming Questions:

Solution 1:

July 11, 2025

```
[5]: import tensorflow as tf
import tensorflow_datasets as tfds
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

dataset, info = tfds.load('tf_flowers', with_info=True, as_supervised=True)
num_classes = info.features['label'].num_classes

train_size = int(0.8 * info.splits['train'].num_examples)
test_size = info.splits['train'].num_examples - train_size

train_ds = dataset['train'].take(train_size)
test_ds = dataset['train'].skip(train_size)

def preprocess(image, label):
    image = tf.image.resize(image, [128, 128])
    image = tf.cast(image, tf.float32) / 255.0
    return image, label

BATCH_SIZE = 32
train_ds = train_ds.map(preprocess).shuffle(1000).batch(BATCH_SIZE).prefetch(tf.data.AUTOTUNE)
test_ds = test_ds.map(preprocess).batch(BATCH_SIZE).prefetch(tf.data.AUTOTUNE)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
```

```

        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dense(num_classes, activation='softmax')
    )

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(train_ds, validation_data=test_ds, epochs=10)

plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Test Accuracy')
plt.title("Training and Testing Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()

y_true = []
y_pred = []

for images, labels in test_ds:
    preds = model.predict(images)
    y_true.extend(labels.numpy())
    y_pred.extend(np.argmax(preds, axis=1))

print("\nClassification Report:")
print(classification_report(y_true, y_pred))

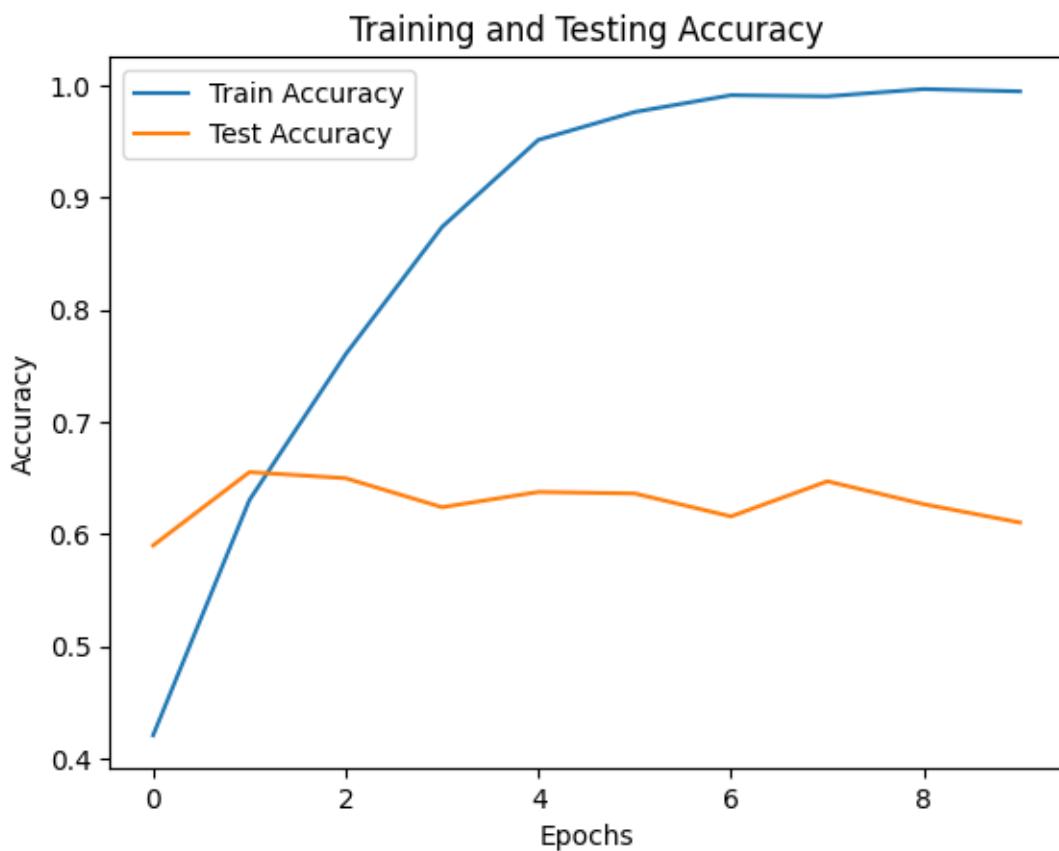
print("Confusion Matrix:")
print(confusion_matrix(y_true, y_pred))

/usr/local/lib/python3.11/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Epoch 1/10
92/92          112s 1s/step -
accuracy: 0.3366 - loss: 2.0018 - val_accuracy: 0.5899 - val_loss: 1.0314
Epoch 2/10

```

```
92/92          93s 1s/step -
accuracy: 0.6191 - loss: 0.9779 - val_accuracy: 0.6553 - val_loss: 0.9203
Epoch 3/10
92/92          95s 1s/step -
accuracy: 0.7318 - loss: 0.7372 - val_accuracy: 0.6499 - val_loss: 0.9015
Epoch 4/10
92/92          136s 964ms/step -
accuracy: 0.8716 - loss: 0.4013 - val_accuracy: 0.6240 - val_loss: 1.1545
Epoch 5/10
92/92          94s 1s/step -
accuracy: 0.9437 - loss: 0.2114 - val_accuracy: 0.6376 - val_loss: 1.4820
Epoch 6/10
92/92          96s 1s/step -
accuracy: 0.9706 - loss: 0.1066 - val_accuracy: 0.6362 - val_loss: 1.3767
Epoch 7/10
92/92          136s 975ms/step -
accuracy: 0.9907 - loss: 0.0468 - val_accuracy: 0.6158 - val_loss: 1.7748
Epoch 8/10
92/92          94s 1s/step -
accuracy: 0.9909 - loss: 0.0500 - val_accuracy: 0.6471 - val_loss: 1.6884
Epoch 9/10
92/92          144s 1s/step -
accuracy: 0.9955 - loss: 0.0210 - val_accuracy: 0.6267 - val_loss: 2.2096
Epoch 10/10
92/92          134s 961ms/step -
accuracy: 0.9967 - loss: 0.0160 - val_accuracy: 0.6104 - val_loss: 1.9694
```



```
1/1      0s 405ms/step
1/1      0s 252ms/step
1/1      0s 242ms/step
1/1      0s 247ms/step
1/1      0s 265ms/step
1/1      0s 244ms/step
1/1      0s 247ms/step
1/1      0s 263ms/step
1/1      0s 245ms/step
1/1      0s 246ms/step
1/1      0s 302ms/step
1/1      0s 258ms/step
1/1      0s 238ms/step
1/1      0s 293ms/step
1/1      0s 270ms/step
1/1      0s 287ms/step
1/1      0s 251ms/step
1/1      0s 248ms/step
1/1      0s 268ms/step
1/1      0s 261ms/step
```

```
1/1          0s 244ms/step  
1/1          0s 263ms/step  
1/1          0s 320ms/step
```

Classification Report:

	precision	recall	f1-score	support
0	0.57	0.77	0.66	158
1	0.75	0.36	0.49	141
2	0.48	0.66	0.56	146
3	0.81	0.69	0.74	151
4	0.61	0.54	0.57	138
accuracy			0.61	734
macro avg	0.64	0.60	0.60	734
weighted avg	0.64	0.61	0.61	734

Confusion Matrix:

```
[[122  7 12 11  6]  
 [ 40 51 30  7 13]  
 [ 14  3 97  4 28]  
 [ 26  1 19 104  1]  
 [ 12  6 43  3 74]]
```

Solution 2:

July 11, 2025

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.model_selection import train_test_split

df = pd.read_csv('A5.csv')
print(df.head())

prices = df.iloc[:, 0].values.reshape(-1, 1)

scaler = MinMaxScaler()
scaled_prices = scaler.fit_transform(prices)

def create_sequences(data, seq_length):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i+seq_length])
        y.append(data[i+seq_length])
    return np.array(X), np.array(y)

seq_len = 60
X, y = create_sequences(scaled_prices, seq_len)

split = int(0.8 * len(X))
X_train, y_train = X[:split], y[:split]
X_test, y_test = X[split:], y[split:]

model = Sequential([
    LSTM(50, return_sequences=True, input_shape=(seq_len, 1)),
```

```

        LSTM(50),
        Dense(1)
    ])

model.compile(optimizer='adam', loss='mean_squared_error')
history = model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))

predictions = model.predict(X_test)
predicted_prices = scaler.inverse_transform(predictions)
real_prices = scaler.inverse_transform(y_test)

plt.figure(figsize=(10,6))
plt.plot(real_prices, label='Actual Price')
plt.plot(predicted_prices, label='Predicted Price')
plt.title('Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()

```

```

      Close
0  24.320436
1  23.635284
2  23.637505
3  23.968964
4  24.889898

/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/rnn.py:200:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
    super().__init__(**kwargs)

Epoch 1/20
49/49          10s 87ms/step -
loss: 0.0171 - val_loss: 0.0022
Epoch 2/20
49/49          5s 80ms/step -
loss: 4.0919e-04 - val_loss: 0.0021
Epoch 3/20
49/49          5s 75ms/step -
loss: 3.5809e-04 - val_loss: 0.0018
Epoch 4/20
49/49          4s 72ms/step -
loss: 3.4392e-04 - val_loss: 0.0018

```

```
Epoch 5/20
49/49          4s 81ms/step -
loss: 3.4887e-04 - val_loss: 0.0021
Epoch 6/20
49/49          5s 91ms/step -
loss: 3.3972e-04 - val_loss: 0.0028
Epoch 7/20
49/49          3s 71ms/step -
loss: 3.5537e-04 - val_loss: 0.0030
Epoch 8/20
49/49          6s 84ms/step -
loss: 3.5901e-04 - val_loss: 0.0019
Epoch 9/20
49/49          5s 75ms/step -
loss: 2.8566e-04 - val_loss: 0.0024
Epoch 10/20
49/49          4s 74ms/step -
loss: 2.7825e-04 - val_loss: 0.0014
Epoch 11/20
49/49          4s 81ms/step -
loss: 2.5916e-04 - val_loss: 0.0020
Epoch 12/20
49/49          5s 75ms/step -
loss: 2.8406e-04 - val_loss: 0.0013
Epoch 13/20
49/49          4s 73ms/step -
loss: 2.5945e-04 - val_loss: 0.0016
Epoch 14/20
49/49          4s 83ms/step -
loss: 2.8621e-04 - val_loss: 0.0014
Epoch 15/20
49/49          4s 89ms/step -
loss: 3.1127e-04 - val_loss: 0.0026
Epoch 16/20
49/49          4s 73ms/step -
loss: 2.6444e-04 - val_loss: 0.0015
Epoch 17/20
49/49          6s 93ms/step -
loss: 2.6091e-04 - val_loss: 0.0017
Epoch 18/20
49/49          4s 83ms/step -
loss: 2.4419e-04 - val_loss: 0.0010
Epoch 19/20
49/49          4s 73ms/step -
loss: 2.9064e-04 - val_loss: 0.0012
Epoch 20/20
49/49          4s 75ms/step -
loss: 2.0154e-04 - val_loss: 0.0015
```

13/13

2s 78ms/step

